



The cost of the agentic approach (RQ3)

Cost per query



Latency (provider pacing removed)



P3 makes ~5 sequential model calls per query versus RAG's one— the source of both the dollar and the latency cost. Mean of five runs.